



EECS 204002

Data Structures 資料結構

Prof. REN-SONG TSAY 蔡仁松 教授

NTHU

CH. 8

HASHING

Motivation

- Operations in a dictionary
 - Get, Insert and Delete
- **Binary search tree**
 - Get, Insert and Delete take $O(n)$
- **Balanced binary search tree (AVL tree)**
 - Get, Insert and Delete take $O(\log n)$
- **Hashing**
 - Get, Insert and Delete take $O(1)$
 - Static hashing
 - Dynamic hashing




8.2

Static Hashing


Overview of Hashing

- The file blocks are divided into M equal-sized *buckets*
- The record with hash key value K is stored in bucket i
 - $i = h(K)$, and h is the *hashing function*




Hash Tables

- **Hash table** (ht) pair=[record, key]
 - A container stores dictionary pairs.
- Hash table is partitioned into **b buckets**
 - $ht[0], ht[1], \dots, ht[b - 1]$
 - Each bucket holds s dictionary pairs (**slots**)
 - Usually $s = 1$, i.e. each bucket can hold exactly one pair.



Hash Function

- The *hash (bucket address)* of a pair with key k is determined by a **hash function**, $h(k)$.
- Hash function maps keys into buckets by returning an integer in the range $[0, b - 1]$.



Definitions

- *Key density (n/T)*
 - n : # of pairs in the table
 - T : Total # of possible keys
- *Loading density or loading factor*
 - $\alpha = n/(s \cdot b)$
- Two keys, k_1 and k_2 , are said to be *synonyms w.r.t. h* , if $h(k_1) = h(k_2)$.

Definitions

- Many keys might be mapped to the same home bucket (*synonyms*)
- *Collision*
 - When a key is mapped to a non-empty home bucket
- *Overflow*
 - When a key is mapped to a full home bucket
- Overflow and collision occur simultaneously when each bucket has 1 slot.

Example

- Given a set of 8 keys ($n = 8$)
 $\{GA, D, A, G, L, A2, A1, A3\}$.
- Consider a h_t with $b = 26$
and $s = 2$.
 - $\alpha = \frac{n}{s \cdot b} = \frac{8}{2 \cdot 26} = 0.154$
- The hash function maps
each key into a bucket using
its leading letter.
 - Represent A – Z as 0 – 25

	slot 1	slot 2
0		
1		
2		
⋮		
⋮		
⋮		
25		

Example (cont'd)

GA, D, A, G, A2 A3 A4

Collision

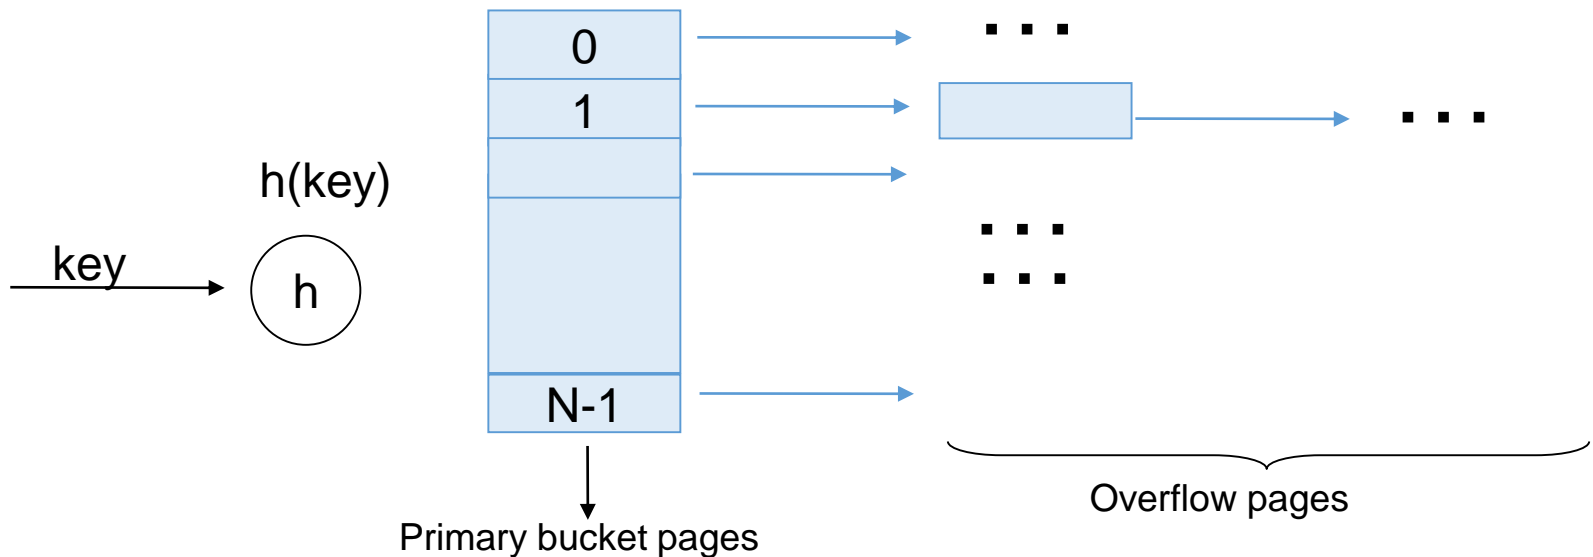
Overflow

mapping	
A	0
D	3
G	6

	slot 1	slot 2	
0	A	A2	A3 A4
1			
2			
3	D		
4			
5			
6	GA	G	
.			
.			
.			
25			

Overflow

- A new record hashes to a bucket that is already full
 - An *overflow file* is kept for storing such records
 - Overflow records that hash to each bucket can be linked together



Hashing Properties

- If the # of slots is small, all operations (search, insert and delete) can be performed in $O(1)$.
- Using leading letter is not a good hash function.
 - Keys might bias toward certain buckets.
- A good hash function should be
 - **Easy to compute**
 - **Few collisions**

Uniform Hash Function

- A hash function that does not result in a biased use of the hash table for **random keys**.
- Given a key k chosen at random, $\text{probability}[h(k) = i] = \frac{1}{b}, \forall i$.
- Four popular hash functions
 - Division
 - Mid-Square
 - Folding
 - Digit Analysis

Division

- $h(k) = k \% D$
- Keys are non-negative integer
- The home bucket is obtained by using the modulo (%) operator.
- Bucket address range from 0 to $D - 1$,
 - hash table must have at least $b = D$ buckets.
- Using a prime number for D (see textbook).
- Ex: $h(k = 219) = 219 \% 8 = 3$

Mid-Square

- Squaring the keys.
- Use an appropriate number of bits from the middle of the squared key as bucket address.
- If r bits is used, the size of the table is 2^r
 - If there are 8 buckets (2^3), we need the middle 3-bits to determine the bucket address

key=219 $219^2=47961=1011$ 1011 0101 1001 $r=3$
 $h(219)=5$

Folding

- The key is partitioned into several parts
- These parts are added together to obtain the key address

$k=12320324111220$

+ + + + = 699

Digit Analysis

- All the keys in the table are known in advance
- Represent each key as a number in radix r
- Digits having the most skewed distributions are deleted
- Employ the remaining digits
- Example: 100 buckets = 2 digits
 - $m = 10^5$ ➔ delete 3 digits

$k_1 =$	d_{11}	d_{12}	...	d_{1n}
$k_2 =$	d_{21}	d_{22}	...	d_{2n}
...				
$k_m =$	d_{m1}	d_{m2}	...	d_{mn}

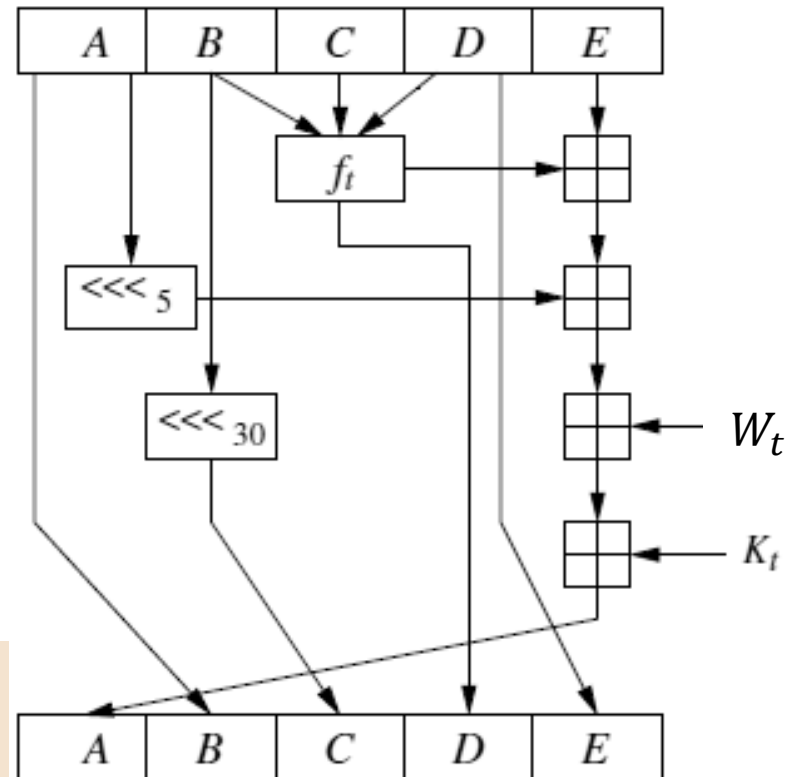
8.2.3

The Secure Hash Algorithm (SHA)

- 4 runs * 20 steps

t	Step #, $0 \leq t \leq 79$
f_t	logical function
\lll_k	Circular left shift k bits
W_t	A 32-bit value derived from M_i
K_t	A constant

$$f_t(B, C, D) = \begin{cases} (B \wedge C) \vee ((\neg B) \wedge D) & \text{if } 0 \leq t \leq 19 \\ B \oplus C \oplus D & \text{if } 20 \leq t \leq 39 \\ (B \wedge C) \vee (B \wedge D) \vee (C \wedge D) & \text{if } 40 \leq t \leq 59 \\ B \oplus C \oplus D & \text{if } 60 \leq t \leq 79 \end{cases}$$



SHA-1: a Merkle-Damgard Hash Function

- Padding: Given an m -bit message, a single bit “1” is appended as the $(m + 1)$ -th bit and then $(448 - (m + 1)) \bmod 512$ (between 0 and 511) zero bits are appended. As a result, the message becomes 64-bit short of being a multiple of 512 bits long.
- Merkle-Damgard Strengthening: append the 64-bit representation of the original length of m , making the result a multiple of 512 bits long.
- Divide the result into 512-bit blocks, denoted by M_1, M_2, \dots, M_l .

m bits	1 bit		64 bits
message	1	000...0	m

SHA-1

- The internal state of SHA-1 is composed of five 32-bit words A, B, C, D and E, used to keep the 160-bit chaining value h_i .
- Initialization: The initial value (h_0) is (in hexadecimal)
 - $A_0 = 67452301x$
 - $B_0 = EFCDAB89x$
 - $C_0 = 98BADCFEx$
 - $D_0 = 10325476x$
 - $E_0 = C3D2E1F0x$.
- Compression: For each block, the compression function $h_i = H(h_{i-1}, M_i)$ is applied on the previous value of $h_{i-1} = (A, B, C, D, E)$ and the message block.
- Output: The hash value is the 160-bit value $h_l = (A, B, C, D, E)$.

The Compression Function H

- Divide M_i into 16*32-bit words:
 - $W_0, W_1, W_2, \dots, W_{15}$.
- for $t = 16$ to 79 compute $W_t = (W_{t-3} \oplus W_{t-8} \oplus W_{t-14} \oplus W_{t-16}) \ll 1$.
 - Remark: The one-bit rotate in computing W_t was not included in SHA, and is the only difference between SHA and SHA-1.
- Set $h_0 = (A_0, B_0, C_0, D_0, E_0)$.
- For $t = 0$ to 79 do
 - $T = A_t \ll 5 + f_t(B_t, C_t, D_t) + E_t + W_t + K_t$.
 - $E_{t+1} = D_t, D_{t+1} = C_t, C_{t+1} = B_t \ll 30, B_{t+1} = A_t, A_{t+1} = T$.
- Output $A = A_0 + A_{80}, B = B_0 + B_{80}, C = C_0 + C_{80}, D = D_0 + D_{80}$, and $E = E_0 + E_{80} \pmod{2^{32}}$.
- The function f_t and the values K_t used above are:

	$f_t(X, Y, Z) =$	$K_t =$
$0 \leq t \leq 19$	$XY \vee (\neg X)Z$	5A827999
$20 \leq t \leq 39$	$X \oplus Y \oplus Z$	6ED9EBA1
$40 \leq t \leq 59$	$XY \vee XZ \vee YZ$	8F1BBCDC
$60 \leq t \leq 79$	$X \oplus Y \oplus Z$	CA62C1D6

Overflow Handling

- Open addressing
 - Linear probing
 - Quadratic probing
 - Rehashing
 - Random probing
- Chaining

Linear Probing: Insert

- Find the closest unfilled bucket.
- To insert a key k .
 - Compute $h(k)$.
 - Check the hash table buckets in the order $h_t[h(k)]$, $h_t[(h(k) + 1) \% b]$, ..., $h_t[(h(k) + j) \% b]$ until an empty bucket is found.
 - If no empty bucket is found, double the size of h_t .
- e.g. GA, D, A, G, A2

0	A
1	A2
2	
3	D
4	
5	GA
6	G
7	
8	
9	
	.
	.
	.

Linear Probing: Search

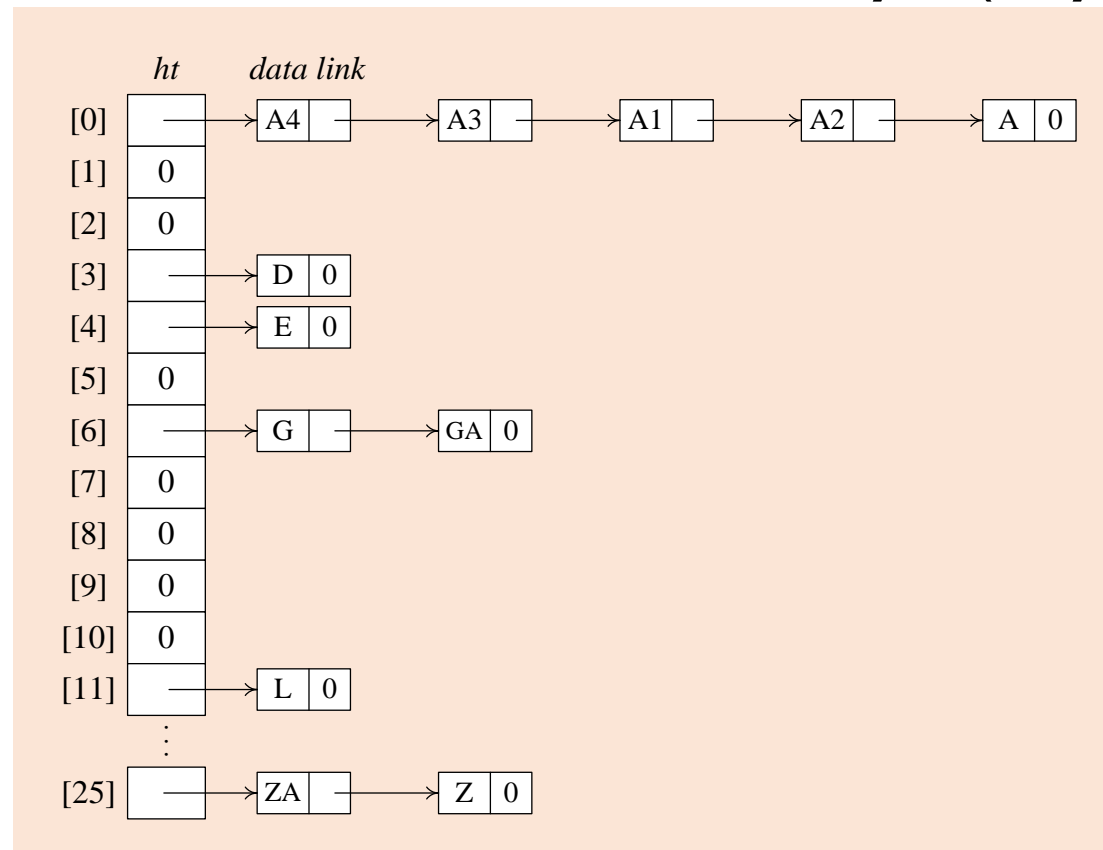
- Searching for a key k .
 - Compute $h(k)$.
 - Examine the hash table buckets in the order $h_t[h(k)], h_t[(h(k) + 1) \% b], \dots, h_t[(h(k) + j) \% b]$ until:
 - $h_t[(h(k) + j) \% b]$ has the same key. Found!
 - $h_t[(h(k) + j) \% b]$ is empty. Not found!
 - Go back to starting point. Not found!
- Disadvantage:
 - Keys tend to cluster together.

Others

- Quadratic probing:
 - Compute $h(k)$.
 - Examine buckets at $h(k)$, $(h(k) + i^2) \% b$, and $(h(k) - i^2) \% b$, $1 \leq i \leq (b - 1)/2$.
- Rehashing:
 - A series of hashing functions h_1, h_2, \dots, h_n .
 - Bucket is searched by h_1, h_2, \dots, h_n .

Chaining

- Use chained hash table to solve collisions
- Each bucket holds a list of keys (key chain)



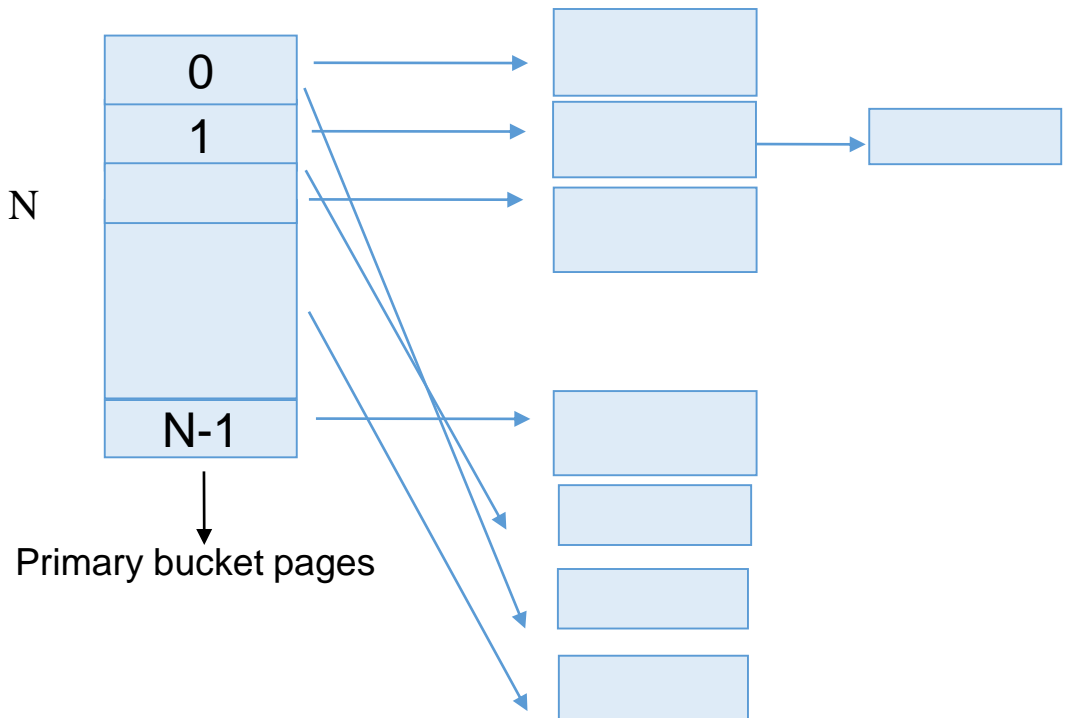



8.3

Dynamic Hashing

Manage Overflow Problems

- Add overflow pages
- Double the size of the buckets
- Double the number of the buckets and reorganize
- ?



**overhead
to rebuild**


Dynamic Hashing

- Also called Extendible Hashing
- Idea: Use directory of pointers to buckets
 - Use the *binary representation* of the hash value $h(K)$ in order to access a *directory*
 - Double #buckets by doubling the directory
 - Splitting just the bucket that is overflowed!
- Directory is much smaller than bucket file
 - Much cheaper to double the directory
 - Split only the page of data entries. *No overflow page!*


Directory

- An array of size 2^d where d is called the *global depth*
- Expand or shrink dynamically
- Entries point to the buckets
 - That contain the stored records
 - When an insertion in a bucket that is full the bucket splits into two buckets
 - The records are redistributed among the two buckets
- Update directory appropriately

Example of Dynamic Hashing

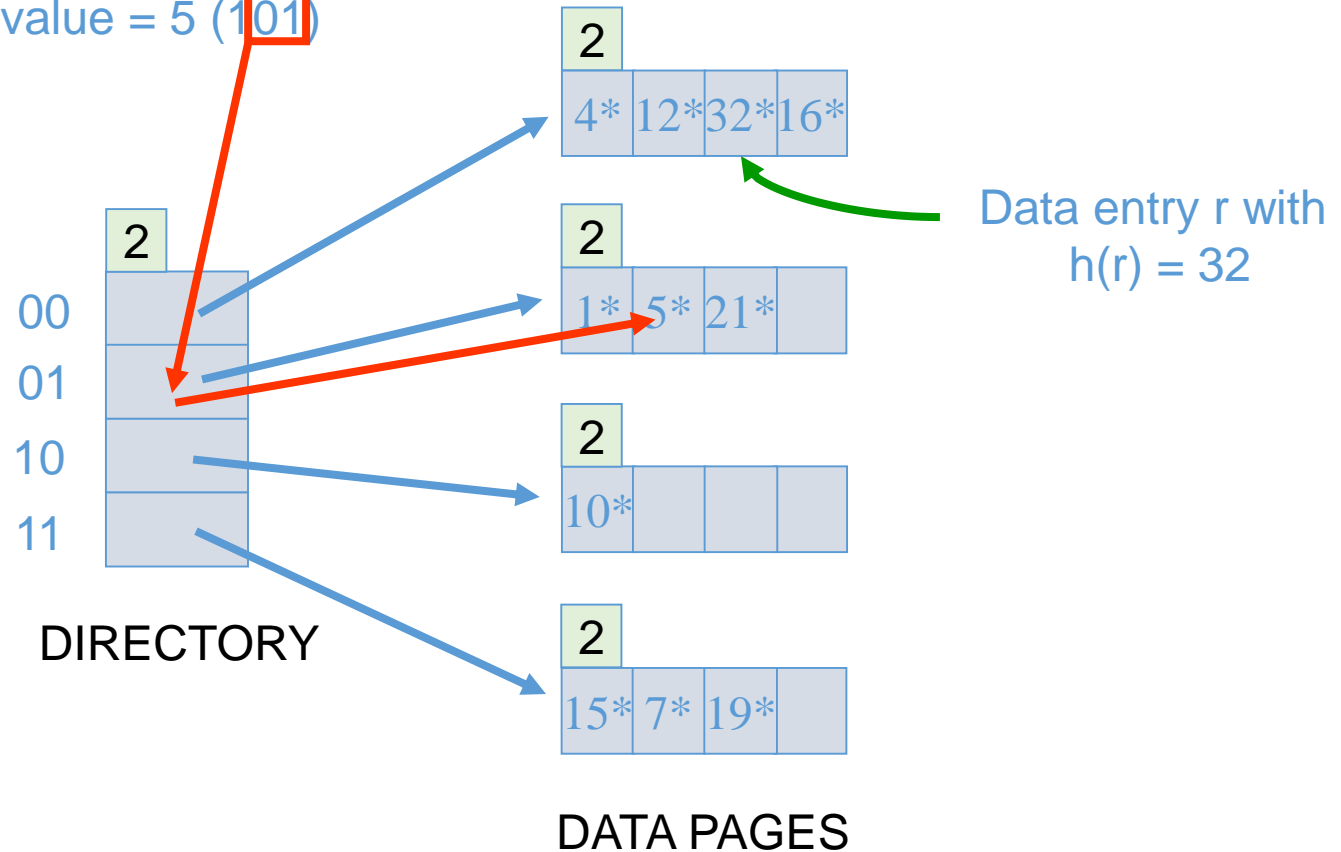


Global Depth = max(Local Depth of all buckets)
Tells # bits needed to determine the address

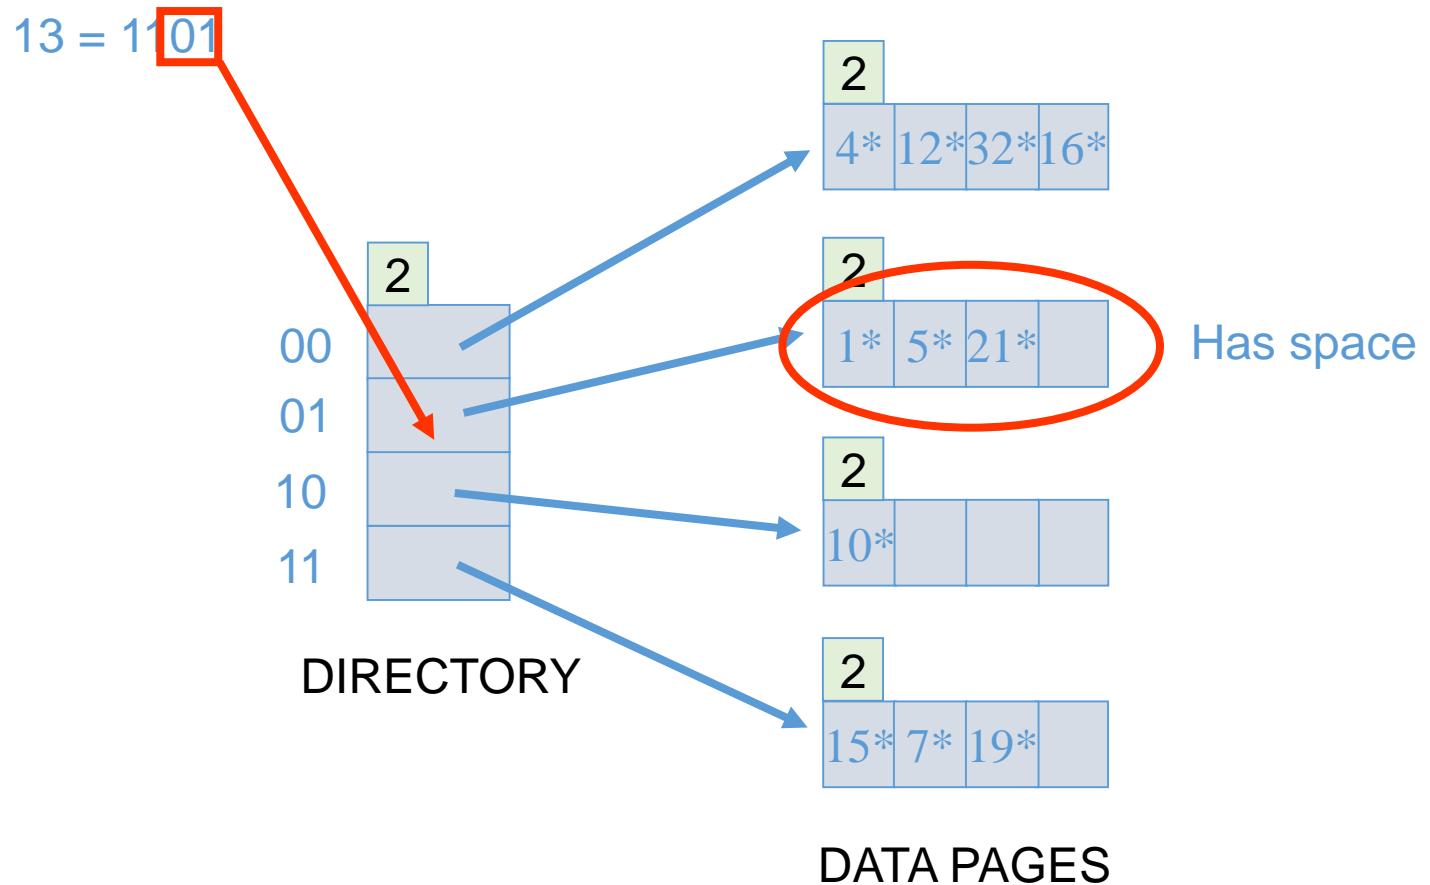


Dynamic Hashing: Example

To locate hash value = 5 (101)




Dynamic Hashing: Insert 13*



Dynamic Hashing: Insert 20*

20 = 10100




Full, need split.
Consider the last
three bits

DIRECTORY


DATA PAGES

Dynamic Hashing: Insert 20*



Dynamic Hashing: Insert 9*


9 = 1001



Need split but directory doubling is not necessary

Dynamic Hashing: Insert 9*

9 = 1001



When to Double Directory

- Initially, all local depths are equal to global depth
 - # of bits need to express the total # of buckets
- During the process of split, if the bucket whose local depth = global depth
 - The directory must be doubled
- Global depth +1 when the directory doubles
 - Local depth +1 when a bucket is split



8.4

Bloom Filters

Introduction

- Generalize the hashing ideas
 - $h(k_1) = h(k_2) \Rightarrow k_1 = k_2$
 - $h(k_1) \neq h(k_2) \Rightarrow k_1 \neq k_2$
- Approximate set membership problem
- Trade-off between the space and the false positive probability

Approximate Set Membership Problem

- Given a set
 $S = \{s_1, s_2, \dots, s_n\} \subseteq U$ (Universe)
- Want to check if “ x is an element of S ”
- Approximated approach
 - $h(S) = h(s_1) \vee h(s_2) \vee \dots \vee h(s_n)$
 - $h(x) \wedge h(S) = \begin{cases} 1 & x \in S & \text{false positive} \\ 0 & x \notin S & \text{sure exclusion} \end{cases}$
 - Take little space

Bloom Filters

1. An n -bit array $A[n]$, initially set to 0
2. k independent random hash functions
$$h_1, \dots, h_k: U \Rightarrow \{0, 1, \dots, n - 1\}$$
3. $\forall s \in S, A[h_i(s)] = 1$ for $1 \leq i \leq k$

To check if $x \in S$, calculate

$$\bigcap_{1 \leq i \leq k} A[h_i(x)] = \begin{cases} 1 & x \in S \\ 0 & x \notin S \end{cases} \quad \begin{array}{l} \text{assume } x \in S \\ \text{sure exclusion} \end{array}$$

BF Design Consideration

- Choose n (filter size in bits).
 - Use as much memory as is available.
- Pick k (number of hash functions).
 - k too small ➡ high probability for different keys to have same signature.
 - k too large ➡ soon to fill ones in the filter
- Select the k hash functions.
 - Hash functions should be relatively independent.

Example: $n = 11, k = 2$

1	0	1	1	0	1	0	1	1	0	0
0	1	2	3	4	5	6	7	8	9	10

- $h_1(s) = s \bmod n$.
- $h_2(s) = 2(s + 1) \bmod n$.

my class
 $m = 3$

[

Student ID	h_1	h_2
105021121	7	5
210510215	3	8
106000103	0	2
107062601	8	7
104034052	1	4

false positive
not in class

The Probability of a False Positive

- We assume the hash function are random.
- After all the elements of S are hashed into the bloom filters, the probability that a specific bit is still 0 is

$$p = (1 - 1/n)^{km} \approx e^{-km/n}$$

Note:

$$e^{-x} = 1 - x + \frac{x^2}{2} - \dots \approx 1 - x$$
$$(e^{-x})^{km} \approx (1 - x)^{km}$$

Optimal n & k

- The probability of a false positive f is
$$f = (1 - p)^k \approx (1 - e^{-km/n})^k$$
- To find the optimal k to minimize f .

Minimize f iff minimize $g = \ln(f)$

$$\frac{dg}{dk} = \ln(1 - e^{-km/n}) + \frac{km}{n} \frac{e^{-\frac{km}{n}}}{(1 - e^{-km/n})}$$

$$\Rightarrow k = \ln(2) * (n/m),$$

$$\Rightarrow f = (1/2)^k = (0.6185)^{n/m}$$

The false positive probability falls exponentially in n/m , the number of bits used per item !!

Conclusion

- A Bloom filters is like a hash table, and simply uses one bit to keep track whether an item hashed to the location.
- If $k = 1$, it's equivalent to a hashing based fingerprint system.
- If $n = cm$ for small constant c , such as $c = 8$, then $k = 5$ or 6 , the false positive probability is just over 2% .
- It's interesting that when k is optimal $k = \ln(2) * (n/m)$, then $p = 1/2$.

An optimized Bloom filters looks like a random bit-string